

In [1]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

In [2]:

```
df=pd.read_csv('Downloads/Heart_Disease_Prediction.csv')
```

In [3]:

```
df.head()
```

Out[3]:

	Age	Sex	Chest pain type	BP	Cholesterol	FBS over 120	EKG results	Max HR	Exercise angina	ST depression	Slope of ST	Number of vessels fluor
0	70	1	4	130	322	0	2	109	0	2.4	2	
1	67	0	3	115	564	0	2	160	0	1.6	2	
2	57	1	2	124	261	0	0	141	0	0.3	1	
3	64	1	4	128	263	0	0	105	1	0.2	2	
4	74	0	2	120	269	0	2	121	1	0.2	1	

In [4]:

```
df.isnull().sum()
```

Out[4]:

```
Age          0
Sex          0
Chest pain type  0
BP           0
Cholesterol  0
FBS over 120 0
EKG results  0
Max HR       0
Exercise angina 0
ST depression 0
Slope of ST  0
Number of vessels fluoro 0
Thallium     0
Heart Disease 0
dtype: int64
```

In [5]:

```
print(df.info())
```

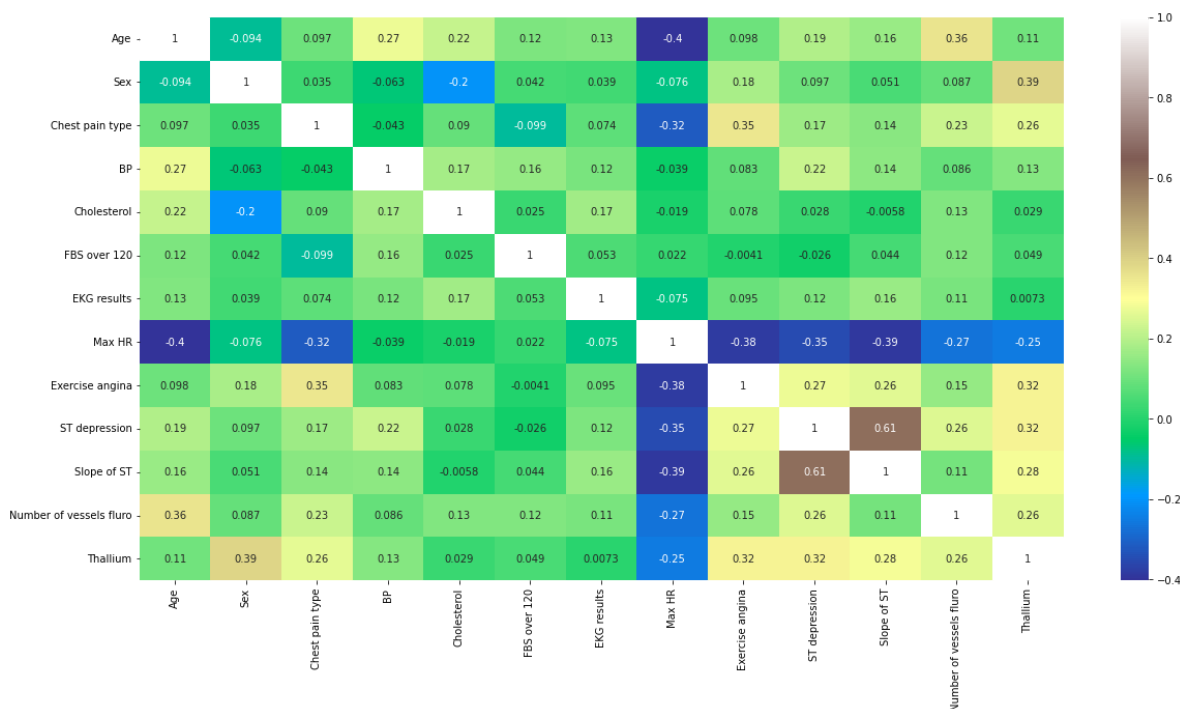
```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 270 entries, 0 to 269
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                    270 non-null    int64
1   Sex                                    270 non-null    int64
2   Chest pain type                        270 non-null    int64
3   BP                                     270 non-null    int64
4   Cholesterol                           270 non-null    int64
5   FBS over 120                          270 non-null    int64
6   EKG results                           270 non-null    int64
7   Max HR                                270 non-null    int64
8   Exercise angina                        270 non-null    int64
9   ST depression                          270 non-null    float64
10  Slope of ST                            270 non-null    int64
11  Number of vessels fluro                270 non-null    int64
12  Thallium                               270 non-null    int64
13  Heart Disease                          270 non-null    object
dtypes: float64(1), int64(12), object(1)
memory usage: 29.7+ KB
None
```

In [6]:

```
plt.figure(figsize=(20,10))
sns.heatmap(df.corr(), annot=True, cmap='terrain')
```

Out[6]:

&lt;AxesSubplot:&gt;



In [7]:

```
sns.pairplot(data=df)
```

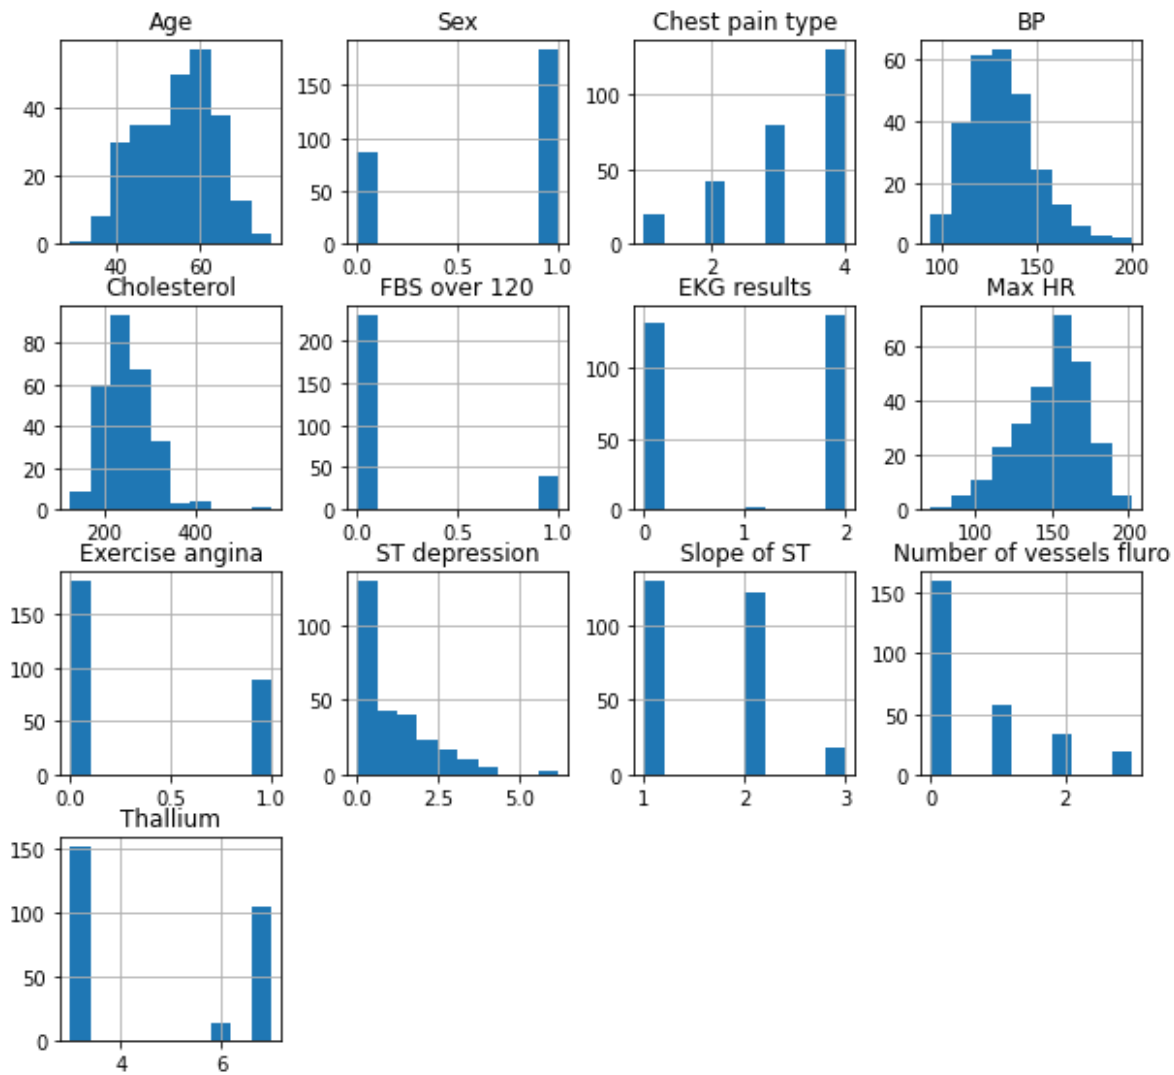
Out[7]:

```
<seaborn.axisgrid.PairGrid at 0x1bb678b6308>
```




In [8]:

```
df.hist(figsize=(10,12), layout=(5,4));
```



In [9]:

```
df.plot(kind='box', subplots=True, layout=(6,3), figsize=(10,10))
plt.show()
```




In [10]:

```
sns.catplot(data=df, x='Sex', y='Age', hue='Heart Disease', palette='tab10')
```

Out[10]:

```
<seaborn.axisgrid.FacetGrid at 0x1bb71a93fc8>
```




In [11]:

```
sns.barplot(data=df, x='Sex', y='Cholesterol', hue='Heart Disease', palette='spring')
```

Out[11]:

```
<AxesSubplot:xlabel='Sex', ylabel='Cholesterol'>
```



In [12]:

```
df['Sex'].value_counts()
```

Out[12]:

```
1    183
0     87
Name: Sex, dtype: int64
```

In [13]:

```
df['Chest pain type'].value_counts()
```

Out[13]:

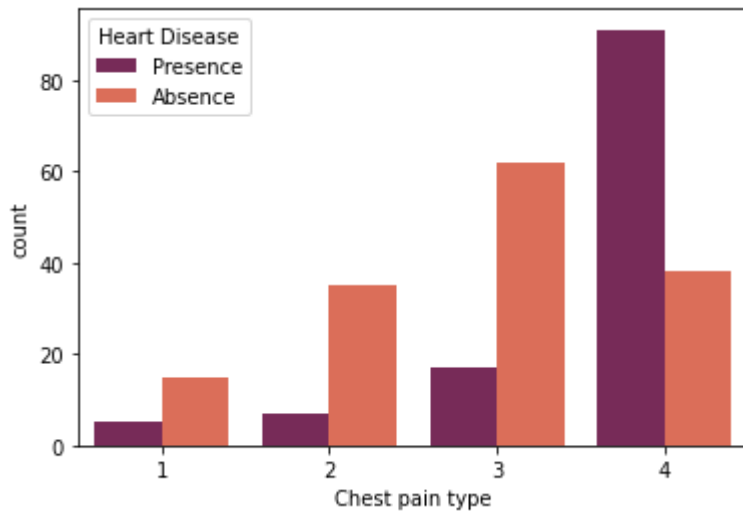
```
4    129
3     79
2     42
1     20
Name: Chest pain type, dtype: int64
```

In [14]:

```
sns.countplot(x='Chest pain type', hue='Heart Disease' , data=df, palette='rocket')
```

Out[14]:

```
<AxesSubplot:xlabel='Chest pain type', ylabel='count'>
```



In [15]:

```
gen = pd.crosstab(df['Sex'], df['Heart Disease'])
print(gen)
```


	Heart Disease	Absence	Presence
Sex			
0		67	20
1		83	100

In [16]:

```
gen.plot(kind='bar', stacked='True', color=['green','blue'],grid=False)
```

Out[16]:

&lt;AxesSubplot:xlabel='Sex'&gt;



In [17]:

```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
StandardScaler = StandardScaler()
columns_to_scale=['Age', 'EKG results', 'Cholesterol', 'Thallium', 'Number of vessels fluro
df[columns_to_scale] = StandardScaler.fit_transform(df[columns_to_scale])
```

In [18]:

```
df.head()
```

Out[18]:

	Age	Sex	Chest pain type	BP	Cholesterol	FBS over 120	EKG results	Max HR	Exercise angina	ST depression	Slope of ST
0	1.712094	1	4	130	1.402212	0	0.981664	109	0	2.4	2
1	1.382140	0	3	115	6.093004	0	0.981664	160	0	1.6	2
2	0.282294	1	2	124	0.219823	0	-1.026285	141	0	0.3	1
3	1.052186	1	4	128	0.258589	0	-1.026285	105	1	0.2	2
4	2.152032	0	2	120	0.374890	0	0.981664	121	1	0.2	1



In [19]:

```

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
StandardScaler = StandardScaler()
columns_to_scale=['Age', 'EKG results', 'Cholesterol', 'Thallium', 'Number of vessels fluro
df[columns_to_scale] = StandardScaler.fit_transform(df[columns_to_scale])

```

In [20]:

```
df.head()
```

Out[20]:

	Age	Sex	Chest pain type	BP	Cholesterol	FBS over 120	EKG results	Max HR	Exercise angina	ST depression	Slope of ST
0	1.712094	1	4	130	1.402212	0	0.981664	109	0	2.4	2
1	1.382140	0	3	115	6.093004	0	0.981664	160	0	1.6	2
2	0.282294	1	2	124	0.219823	0	-1.026285	141	0	0.3	1
3	1.052186	1	4	128	0.258589	0	-1.026285	105	1	0.2	2
4	2.152032	0	2	120	0.374890	0	0.981664	121	1	0.2	1

In [21]:

```

x=df.drop(['Heart Disease'], axis=1)
y=df['Heart Disease']

```

In [22]:

```
x_train, x_test, y_train, y_test=train_test_split(x,y,test_size=0.3, random_state=40)
```

In [23]:

```

print('x_train-', x_train.size)
print('x_test-', x_test.size)
print('y_train-', y_train.size)
print('x_test-', x_test.size)

```

```

x_train- 2457
x_test- 1053
y_train- 189
x_test- 1053

```

In [24]:

```
from sklearn.linear_model import LogisticRegression
lr=LogisticRegression()
model1=lr.fit(x_train,y_train)
prediction1=model1.predict(x_test)
```

In [25]:

```
from sklearn.metrics import confusion_matrix
cm=confusion_matrix(y_test,prediction1)
cm
```

Out[25]:


```
array([[40,  5],
       [ 9, 27]], dtype=int64)
```

In [26]:

```
sns.heatmap(cm, annot=True,cmap='BuPu')
```

Out[26]:

&lt;AxesSubplot:&gt;



In [27]:

```
TP=cm[0][0]
TN=cm[1][1]
FN=cm[1][0]
FP=cm[0][1]
print('Testing Accuracy:', (TP+TN+FN)/(TP+TN+FN+FP))
```

Testing Accuracy: 0.9382716049382716

In [28]:

```
from sklearn.metrics import accuracy_score
accuracy_score(y_test,prediction1)
l=accuracy_score(y_test,prediction1)
```

In [29]:

```
from sklearn.metrics import classification_report
print(classification_report(y_test, prediction1))
```

	precision	recall	f1-score	support
Absence	0.82	0.89	0.85	45
Presence	0.84	0.75	0.79	36
accuracy			0.83	81
macro avg	0.83	0.82	0.82	81
weighted avg	0.83	0.83	0.83	81

In [30]:

```
import pandas as pd
from sklearn import neighbors, metrics
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
import numpy as np
import pickle
from sklearn.ensemble import RandomForestClassifier
import pandas as pd
from sklearn.tree import DecisionTreeClassifier
import seaborn as sns
import matplotlib.pyplot as plt
```

In [31]:

```
from sklearn.metrics import accuracy_score
```

In [32]:

```
dataset = pd.read_csv("Downloads/Heart_Disease_Prediction.csv")
```

In [33]:

```
KX = dataset[['Age', 'Sex', 'Chest pain type', 'BP', 'Cholesterol', 'FBS over 120', 'EKG results']
```

In [34]:

```
KY = dataset[['Heart Disease']].values
```

In [35]:

KX

Out[35]:

```
array([[70., 1., 4., ..., 2., 3., 3.],
       [67., 0., 3., ..., 2., 0., 7.],
       [57., 1., 2., ..., 1., 0., 7.],
       ...,
       [56., 0., 2., ..., 2., 0., 3.],
       [57., 1., 4., ..., 2., 0., 6.],
       [67., 1., 4., ..., 2., 3., 3.]])
```

In [36]:

```
KY = KY.flatten()
print(KY)
```

[illegible]

In [37]:

```
KX_train , KX_test , KY_train , KY_test = train_test_split(KX,KY,test_size=0.2,random_state
```

In [38]:

```
knn = KNeighborsClassifier(n_neighbors = 20)
knn.fit(KX_train, KY_train)
print(knn.score(KX_test, KY_test))
```

0.6111111111111112

In [39]:

```
pickle.dump(knn,open('heart_knn_model.sav','wb'))
```

In [40]:

```
predict_knn = knn.predict(KX_test)
accuracy_knn = metrics.accuracy_score(KY_test,predict_knn)
```

In [41]:

```
predict_knn
```

Out[41]:

```
array(['Absence', 'Absence', 'Absence', 'Presence', 'Presence', 'Absence',
       'Presence', 'Absence', 'Absence', 'Absence', 'Presence', 'Absence',
       'Absence', 'Presence', 'Presence', 'Absence', 'Absence', 'Absence',
       'Presence', 'Presence', 'Presence', 'Presence', 'Absence',
       'Absence', 'Absence', 'Presence', 'Presence', 'Absence', 'Absence',
       'Presence', 'Absence', 'Presence', 'Presence', 'Absence',
       'Absence', 'Absence', 'Absence', 'Absence', 'Absence', 'Presence',
       'Absence', 'Presence', 'Absence', 'Absence', 'Absence', 'Absence',
       'Presence', 'Absence', 'Absence', 'Absence', 'Presence', 'Absence', 'Absence',
       'Absence', 'Absence'], dtype=object)
```

In [42]:

```
accuracy_knn
```

Out[42]:

0.6111111111111112

In [43]:

```
k=accuracy_knn
```

In [45]:

```
import csv
import pandas as pd
import numpy as np
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split
from sklearn import metrics
from sklearn.metrics import confusion_matrix, f1_score, roc_curve, auc
import matplotlib.pyplot as plt
from itertools import cycle
from scipy import interp
```

In [46]:

```
df = pd.read_csv('Downloads/Heart_Disease_Prediction.csv', header = None)
```

In [47]:

```
training_x=df.iloc[1:df.shape[0],0:13]
```

In [48]:

```
training_y=df.iloc[1:df.shape[0],13:14]
```

In [49]:

```
nx=np.array(training_x)
ny=np.array(training_y)
```

In [52]:

```
for z in range(5):
    print("\nTest Train Split no. ",z+1,"\n")
    nx_train,nx_test,ny_train,ny_test = train_test_split(nx,ny,test_size=0.25,random_state=
    # Gaussian function of sklearn
    gnb = GaussianNB()
    gnb.fit(nx_train, ny_train.ravel())
    ny_pred = gnb.predict(nx_test)
```

Test Train Split no. 1

Test Train Split no. 2

Test Train Split no. 3

Test Train Split no. 4

Test Train Split no. 5

In [61]:

```
print("\n Naive Bayes model accuracy(in %):", metrics.accuracy_score(ny_test, ny_pred))
```

Naive Bayes model accuracy(in %): 0.7794117647058824

In [62]:

```
n=metrics.accuracy_score(ny_test, ny_pred)
```

In [64]:

```
import pandas as pd
from sklearn import neighbors, metrics
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
import numpy as np
import pickle
from sklearn.ensemble import RandomForestClassifier
import pandas as pd
from sklearn.tree import DecisionTreeClassifier
import seaborn as sns
import matplotlib.pyplot as plt
```

In [65]:

```
from sklearn.metrics import accuracy_score
```

In [67]:

```
dataset = pd.read_csv("Downloads/Heart_Disease_Prediction.csv")
```

In [69]:

```
DX = dataset[['Age', 'Sex', 'Chest pain type', 'BP', 'Cholesterol', 'FBS over 120', 'EKG results']
```

In [70]:

```
dy = dataset[['Heart Disease']].values
```

In [71]:

```
DX
```

Out[71]:

```
array([[70., 1., 4., ..., 2., 3., 3.],
       [67., 0., 3., ..., 2., 0., 7.],
       [57., 1., 2., ..., 1., 0., 7.],
       ...,
       [56., 0., 2., ..., 2., 0., 3.],
       [57., 1., 4., ..., 2., 0., 6.],
       [67., 1., 4., ..., 2., 3., 3.]])
```

In [72]:

```
dy = dy.flatten()
print(dy)
```

[illegible]

In [73]:

```
DX train , DX test , dy train , dy test = train test split(DX,dy,test size=0.2,random state
```

In [74]:

```
from sklearn.tree import DecisionTreeClassifier

max_accuracy = 0
```



In [75]:

```

for x in range(200):
    dt = DecisionTreeClassifier(random_state=x)
    dt.fit(DX_train,dy_train)
    dy_pred_dt = dt.predict(DX_test)
    current_accuracy = round(accuracy_score(dy_pred_dt,dy_test)*100,2)
    if(current_accuracy>max_accuracy):
        max_accuracy = current_accuracy
        best_x = x

```

In [85]:

```

dt = DecisionTreeClassifier(random_state=best_x)
dt.fit(DX_train,dy_train)
dy_pred_dt = dt.predict(DX_test)

```

In [88]:

```
score_dt = (accuracy_score(dy_pred_dt,dy_test))
```

In [89]:

```
print("The accuracy score achieved using Decision Tree is: "+str(score_dt))
```

The accuracy score achieved using Decision Tree is: 0.7962962962962963

In [90]:

```
d=(accuracy_score(dy_pred_dt,dy_test))
```

In [91]:

```

print('Logistic Regression :',l)
print('KNN :',k)
print('Naive Bayes :',n)
print('Decision Tree :',d)

```

Logistic Regression : 0.8271604938271605  
 KNN : 0.6111111111111112  
 Naive Bayes : 0.7794117647058824  
 Decision Tree : 0.7962962962962963

In [93]:

```

print('Logistic Regression :',l*100,'%')
print('KNN :',k*100,'%')
print('Naive Bayes :',n*100,'%')
print('Decision Tree :',d*100,'%')

```

Logistic Regression : 82.71604938271605 %  
 KNN : 61.111111111111114 %  
 Naive Bayes : 77.94117647058823 %  
 Decision Tree : 79.62962962962963 %

In [ ]: